

PROFESSIONAL SUMMARY

GenAI Full-Stack Engineer with 2.6+ years of production experience building enterprise applications and shipping AI-powered products end-to-end. Designed and deployed a RAG-based Knowledge Assistant (Llama 3.3 70B, semantic retrieval, ~1s end-to-end latency) and JobFit, a production AI SaaS reaching 200+ users and generating live revenue. Strong foundation in full-stack engineering (React, Next.js, .NET, PostgreSQL) from 2.6 years at HCL Tech, now focused on applying that foundation to LLM-powered products — from embeddings and vector search to grounded generation and agentic workflows.

SKILLS

Technical Skills:	Generative AI, RAG, Python, Embeddings, Semantic Search, Vector Search, Cosine Similarity, LLM APIs, Grounded Generation, React, Next.js, Angular, TypeScript, JavaScript, Tailwind CSS, Node.js, Express, FastAPI, C#, .NET, REST APIs, Microservices, Authentication, Authorization, Caching, SQL, PostgreSQL, MongoDB, Prisma, Vector Storage, Embedding Indexing
Tools & Environment:	Git, Docker, VS Code, Postman, Vercel, Render, GitHub

PROFESSIONAL EXPERIENCE

Full Stack Developer at HCL Tech	2024-03 - Present
- Developed and integrated 100+ REST APIs that powered enterprise manufacturing workflows, reinforcing API centric design and seamless data exchange between frontend, backend, and data access layers. - Optimized application and API performance, reducing average response latency from ~5/seconds to ~1/second and achieving a five fold speed improvement through code level refinements and efficient data handling. - Delivered 50+ enterprise features across 10 application modules for a manufacturing software platform, driving full stack software development using Angular frontend, .NET APIs, and database layers within a 7 member engineering team. - Resolved 80+ production issues across UI, APIs, and business logic by conducting root cause analysis, cross layer debugging, and delivering production ready fixes.	

KEY PROJECTS

AI Knowledge Assistant – RAG System
Built an end-to-end document grounded AI Knowledge Assistant supporting authenticated document upload, conversational querying, and streamed responses. Designed the RAG ingestion pipeline using 400 token chunks with 40 token overlap and 384 dimensional all MiniLM L6 v2 embeddings across 30+ test documents. Implemented semantic retrieval using cosine similarity, Top K=5, and a 0.3 similarity threshold to select relevant context before LLM generation. Integrated Llama 3.3 70B via Groq with temperature 0.2, 512 token output control, and a 30 second timeout; measured approximately 500ms retrieval and ~1 second end to end latency across 50 test queries. Implemented authenticated document/chat workflows, protected API routes, authorization, streaming responses, and centralized AI configuration for retrieval and generation parameters.
JobFit – ATS Resume Optimization Platform
Built and deployed a production AI SaaS platform reaching 200+ users, 300+ resumes analyzed, and 300+ job descriptions analyzed. Developed AI assisted resume/JD matching with job specific keyword extraction, multi dimensional scoring, and a 0–100 job fit score. Implemented PDF/DOCX processing, structured resume extraction, AI generated improvements, ATS friendly document generation, authentication, and PostgreSQL/Prisma data models. Integrated Razorpay for monetization; reached 20+ paid users and ~13,000 revenue. Generated 6.22K Google impressions, 303 clicks, 4.9% CTR, and 9.7 average search position during the measured period.

EDUCATION

Bachelor of Technology (Information Technology) — Sri Venkateswara College of Engineering